

TracerX: Enhancing Dynamic Symbolic Execution with Weakest Precondition Interpolation

Arpita Dutta¹, Rasool Maghareh², Joxan Jaffar³, Sangharatna Godbole⁴, and Xiao Liang Yu⁴

School of Computing, National University of Singapore, Singapore^{1,3,5}, Lemurian Labs, Toronto², National Institute of Technology Warangal, India⁴

Test-Comp 2025,
[8 May 2025, Hamilton, Canada]



From KLEE To TracerX

- DFS Forward Symbolic Execution to find feasible paths (Similar to KLEE)
- Intermediate execution states preserved (Unlike KLEE)
- Path interpolants are generated for each path during backward tracking
- Tree interpolants are generated as conjunction of path interpolants
- Tree interpolants then used for subsumption at similar program points





Figure: TracerX Framework



Figure: Pruning of subtree

Symbolic Execution Tree with Interpolation



```
x = 0;  
if (b1) x += 3 else x += 2  
if (b2) x += 5 else x += 7  
if (b3) x += 9 else x += 14  
assert(x <= 24)
```

- **Without interpolation:** The full tree is traversed.

- **With interpolation:**

- 1 $\langle 8b \rangle$ context contains $x = 10$. It is subsumed with the tree interpolant from $\langle 8a \rangle$: $x \leq 10$.
- 2 $\langle 5b \rangle$ context contains $x = 2$. Subsumed with the tree interpolant from $\langle 5a \rangle$: $x \leq 3$.
- 3 Big subtree traversal is avoided.

TracerX implementation of Path Based Weakest Precondition (PBWP)

- **Ideal interpolant** is the weakest precondition of the target. Unfortunately, PBWP is **intractable** to compute.
- For example, Assume $(b1 \wedge \neg b2 \wedge \neg b3)$ is **UNSAT**.
WP before first “if-statement” is: $b1 \longrightarrow (\neg b2 \wedge b3 \wedge x \leq 7) \vee (b2 \wedge x \leq 4)$
 $\neg b1 \longrightarrow x < 3$
- Essentially, PBWP is **exponentially disjunctive**
- Challenge is to obtain a conjunctive approximation

A Path is a sequence of assignment and assume instructions:

- 1 Interpolant of **Assignment** instruction:
 - $WP(inst, \omega) = \dots$ inverse transition of $inst$ over ω
 - e.g. $\omega : x \leq 15$ and $inst : x = z + 2$, then $WP(inst, \omega) : z \leq 13$
- 2 Interpolant of **Assume** instruction (C is incoming Context): $\{C\} \text{ assume}(B) \{\omega\}$
 - PBWP Approximation: find \bar{C} to replace C
 - ABDUCTION PROBLEM !!!

Approximation of Path Based Weakest Precondition

This algorithm is the **heart of TracerX**:

- 1 We compute finest partition so that $\text{var}(C_i) * \text{var}(C_j)$ s.t. $i \neq j$:
 $\{C_1 * C_2 * C_3 * \dots * C_n\}$ $\text{assume}(B)$ $\{\omega_1 * \omega_2 * \omega_3 * \dots * \omega_m\}$ (* is as in separation logic).
- 2 Bunch C_i into three:

- **Target independent:** The C_i which are separate from B and ω .
Action: Replace C_i with *true*, i.e. remove C_i .
- **Guard independent:** Consider $C_{gi} \equiv C_i$ s.t. $C_i * B$; and, $\omega_{gi} \equiv \omega_j$ s.t. $B * \omega_j$.
Action: Replace C_{gi} by ω_{gi} .
- **Remainder of the C_i :** We do not capture exact WP for this group.
e.g. $\{z == 5\}$ $\text{assume}(x > z - 2)$ $\{x > 0\}$ (Here, $z > 2$ is the WP)
Action: No change to C_i , i.e. keep C_i .

Challenge

Verification of looping programs up to the maximum (including unbounded loops) bound.

Limitations of Existing Techniques

- Classic symbolic execution (e.g., KLEE[1]) does not support pruning.
- Static symbolic execution (e.g., CBMC[5]) requires the loop unrolling bound upfront.
- Iterative deepening is not incremental. Not recognizing the same program points on loop unroll.
- Interpolation techniques available for model checking are ineffective towards inductive reasoning (eg., CPAchecker[8]).

Our Ultimate Goal

Given a looping program:

- 1 Prove the program for a given loop bound.
- 2 Prove the program for as large as possible loop bound.
- 3 Prove the program for unbounded iterations.

Harness to Explain Incremental Behavior

```
#define loop for(;;)
#define MAXDEEP ∞
<1> main() {
<2>     int B = 1, k, x = 0; /* Initilization of variables */
<3>     loop {
<4>         if ( * ) { /* Non deterministic choice */
<5>             if (B == MAXDEEP) exit(0); /* Fixed Bound Verification */
<6>             B++;
<7>         } else {
<8>             k = B;
<9>             while (k) {
<10>                 x =  $\Psi(x)$ ; /* Program under Test */
<11>                 k = k - 1;
<12>             }
<13>             assert( $\Phi(x)$ ); /* Safety Condition Check */
<14>             break;
<15>         } //EndIf
<16>     } //EndLoop
<17> END;;
<18> }
```

- In line <10>, $\Psi(x)$ is your program under test. It could contain non-determinism but must be free from unbounded loops.
- B and k are ghost variables.
- Line <5> check is used for fixed bound verification.

Framework behind Incremental Analysis of Loops




Figure: Incremental Deepening using Interpolation

For an unbounded program P:

We perform below two sequential checks to find the fixed point:

- First, confirm the tree explored for the most recent bound is subsumption closed.
→ i.e., No new path explored.
- Secondly, the safety condition at bound k is not weaker than at bound $k-1$.
→ We ensure this by equivalence check between the interpolants at bounds $k-1$ and k . $[inp_{k-1} \equiv inp_k]$

Example 1: Incremental Reasoning

```

<1> int B = 1, x = 0;
<2> for (;;) {
<3>     if (*) B++;
<4>     else {
<5>         k = B;
<6>         while (k) {
<7>             x = x + 2;
<8>             k = k - 1;
<9>         }
<10>         assert(x != 99999);
<11>     }
<12> }
```

- TracerX verifies the program up to **10106 iterations** in 3600 seconds.
- CBMC results in **Timeout** for bound $\sim \leq 300$ with a limit of 3600 seconds.
- TracerX successfully utilizes the incremental behavior to reach a higher loop bound.
- TracerX unable to obtain a FixPoint but it can find FixPoint if assume $(x \geq 0)$.

Example 2: Inductive Reasoning when Assertion is Inductive

```
(1) int x = 0;  
(2) int i = 0;  
(3) while (1) {  
(4)     x += 3;  
(5)     if (i == 3) x++;  
(6)     i++;  
(7)     if (i == 2) i = 0;  
(8) }  
(9) assert(x % 3 == 0)
```

Table: Comparison of Interpolation Algorithms for FixPoint Reasoning

Method	Fixed Point	#Unrolling	Invariant	Time
IMC [9]	No	41	-	900s
ISMC [8]	No	16	-	900s
DAR [8]	No	5	-	900s
TracerX	Yes	3	$i=0 \wedge x\%3=0$	0.07s

Example 3: Inductive Reasoning when Assertion is Not Inductive

```

<1> int B = 1, x = 1, y = 1;
<2> for (;;) {
<3>     if (*) B++;
<4>     else {
<5>         k = B;
<6>         while (k) {
<7>             int temp = x;
<8>             x = x + y;
<9>             y = y + temp;
<10>            k = k - 1;
<11>        }
<12>        assert(y >= 1);
<13>    }
<14> }
```

- This example is taken from [10].
- Here, the invariant is not inductive.
- TracerX obtained the fixed point in 2 iterations.
- Obtained inductive interpolant is $(x + y \geq 1)$.

Example 4: Inductive Parametric Reasoning

```
(1) int B = 1, sum = 0;
(2) for (;;) {
(3)     if(*) B++;
(4)     else {
(5)         k = B;
(6)         while (k) {
(7)             sum = 2 + sum;
(8)             k = k - 1;
(9)         }
(10)        assert(sum = 2*B);
(11)    }
(12) }
```

- TracerX obtained the fixed point in 2 iterations.
- Obtained interpolants:
 - At $k==1$: INP_1 is $(sum = (-2 + 2*B))$
 - At $k==2$: INP_2 is $(sum = (-4 + 2*B))$
- On replacing, B by B+1 (since, step-size for B is 1) in INP_2 , : $INP_1 \equiv INP_2$.
- This is true for every k^{th} and $k+1^{th}$ interpolants.

Ex 5: Inductive Parametric Reasoning when Assertion is Not Inductive

```
(1) int B = 1, x1 = 0, x2 = 1;
(2) for (;;) {
(3)     if(*) B++;
(4)     else {
(5)         k = B;
(6)         while (k) {
(7)             int temp = x1;
(8)             x1 = x2;
(9)             x2 = x1 + temp;
(10)            k = k - 1;
(11)        }
(12)        assert(fib(B) == x2);
(13)    }
(14) }
```

- This is the Fibonacci sequence program.
- Function `fib(B)` returns the B^{th} number from Fibonacci sequence.
- Here, the parametric loop invariant is itself not inductive.
- Work in progress.

Experimental Results

Data set: All C-programs (with at least 1 bug) from **RERS Challenge [2012-2022]** [6].

- Removed Programs with more than 300K LOC.
- Total Programs: 137 - 19 (Large) - 6 (From 2022) = **112**
- All programs are event-condition-action (ECA) type with unbounded loops.

Table: Program Characteristics

Year	# Programs	LOC		# Predicates		Size of Predicates		# Target
		Min	Max	Min	Max	Min	Max	
2012	19	595	184.9K	127	17917	2	30	61
2013	18	2.4K	153.3K	279	11152	4	12	60
2014	24	1K	285K	134	22671	2	19	100
2016	14	1.8K	155.7K	199	2862	2	20	100
2017	9	1.9K	140.8K	216	2883	2	20	100
2018	9	1.2K	114.2K	211	2742	2	21	100
2019	11	1.1K	88.6K	200	2702	2	20	100
2020	8	1.1K	127K	199	2825	2	20	100
2022	6	3.2K	270.2K	-	-	-	-	-

- Removed **shallow bugs** using KLEE[1] for 3600 seconds.

Experimental Results

- Both the systems **CBMC** [5] and **TracerX** are run for 3600 seconds
- We first run TracerX to determine the loop bound for CBMC.
- TracerX(Win)**: Programs not finished by CBMC.
- CBMC(Win)**: CBMC finished faster than TracerX.

Table: Comparison of TracerX with CBMC

RERS Year	Total Programs	TracerX (Win)	CBMC (Win)	DeepBug Detected*	Bound Covered	
					TracerX	CBMC
					[Min - Max]	[Min - Max]
2012	19	14	5	4	[4 - 523]	[8 - 377]
2013	18	14	4	1	[4 - 15]	[5 - 15]
2014	24	15	9	0	[4 - 20]	[4 - 20]
2016	14	9	5	3	[5 - 16]	[9 - 16]
2017	9	6	3	1	[3 - 11]	[9 - 11]
2018	9	5	3	1	[4 - 15]	[4 - 11]
2019	11	8	3	2	[3 - 11]	[10 - 12]
2020	8	4	3	0	[4 - 12]	[10-12]

- ① **Website:** <https://tracer-x.github.io/>
- ② **Github:** <https://github.com/tracer-x/>
- ③ **TracerX: Dynamic Symbolic Execution with Interpolation**
J. Jaffar, R. Maghareh, S. Godbole, X.L. Ha, 2020
<https://arxiv.org/abs/2012.00556>
- ④ **TracerX: Dynamic Symbolic Execution with Interpolation (competition contribution)**
J. Jaffar, R. Maghareh, S. Godbole, X.L. Ha, *FASE 2020*
- ⑤ **Toward Optimal MC/DC Test Case Generation**
S. Godbole, J. Jaffar, R. Maghareh, A. Dutta, *ISSTA 2021*
- ⑥ **TracerX: Pruning Dynamic Symbolic Execution with Deletion and Weakest Precondition Interpolation (competition contribution)**
A. Dutta, R. Maghareh, J. Jaffar, S. Godbole, X. L. Yu, *FASE 2024*

- [1] C. Cadar et al. Klee: Unassisted and automatic generation of high-coverage tests for complex systems programs. In: OSDI, 2008.
- [2] J. Jaffar et al. TRACER: A symbolic execution tool for verification. In: CAV, 2012.
- [3] J. Jaffar et al. TracerX: Dynamic symbolic execution with interpolation (competition contribution) . In: FASE, 2020.
- [4] A. Dutta et al. TracerX: Pruning Dynamic Symbolic Execution with Deletion and Weakest Precondition Interpolation (competition contribution). In: FASE, 2024.
- [5] D. Kroening et al. CBMC-C Bounded Model Checker. In: TACAS 2014.
- [6] <http://rers-challenge.org/>
- [7] D. Beyer et al. Augmenting interpolation-based model checking with auxiliary invariants. In International Symposium on Model Checking Software, 2024.
- [8] D. Beyer et al. A transferability study of interpolation-based hardware model checking for software verification. In: FSE, 2024.
- [9] D. Beyer et al. Interpolation and SAT-based model checking revisited: Adoption to software verification. In: J. Autom. Reasoning, 2024.
- [10] A. R. Bradley. Understanding IC3. In: International Conference on Theory and Applications of Satisfiability Testing, 2012.